

CPP Programming: Session 8

Anupama Chandrasekhar (September 27 2018)

STL : The Standard Template Library

- The STL was invented by Alex Stepanov and it offers an efficient way to store, access and view data and is designed to scale from few to millions of elements.

```
vector<int> myVector(NUM_INTS);  
sort(myVector.begin(), myVector.end());
```

- If we did not have an easy way to represent a list of objects, like the STL container **vector**, we would have to deal with memory allocation and freeing that is needed while allocating dynamic memory ... not fun!

Overview of STL

- STL is logically divided into six pieces:
 - **Containers:** E.g.: list of elements can be stored in an vector, associative collection of key/value pairs in a map etc.
 - **Iterators:** To traverse the containers
 - **Algorithms:** STL algorithms operator over ranges of data provided by the iterators, there are algorithms for searching, sorting, reordering, permuting etc.
 - **Adapters:** ADTs built on containers
 - **Functors :** Customizations
 - **Allocators:** Customizations

Why STL? Motivational Example

- When you know exactly how much data your program needs, we can use static allocations. But often that is not the case.
- Take a simple example:
 - Write a program that reads three integers from a user and prints them out in sorted order.
 - Problems we need to solve:
 - How would we store the numbers that the user enters?
 - How would we sort them?
 - See `1_sort3numbers.cpp`
 - Another problem is scalability .. What if we had to sort four numbers.

STL : **vector**

- Fortunately, the STL provides us with a versatile tool called the **vector** that allows us to store sequences of elements using a single variable.
- **vector** can be used to store a sequence of elements of the same type.

Value	137	42	2718	3141	410
Index	0	1	2	3	4

- That is you can have a **vector** of ints, **vector** of **strings** or a **vector** of your own type.

```
//Declaration  
vector<int> int_vector;
```

STL: **vector**

- Using a **vector**. See 2_vector.cpp
- Initially empty, you can use **push_back** to put elements in.
- For sorting, let's implement a simple sorting algorithm
 - Find the smallest element in a list
 - Put that element in the front of the list
 - Repeat until all elements are in place
- Things to note:
 - Always define the template parameter while declaring a **vector**: **vector<int>**, **vector<string>** etc.
 - Prefer to pass by reference for efficiency
 - Use **size_t** for indexing as an index is always positive and **size_t** can hold the largest size.

STL: *vector*

- Another way to do selection sort

Value	100	200	300	400
Index	0	1	2	3

Value	100	137	200	300	400
Index	0	1	2	3	4

- Every time we get a new element insert it in the sorted position and maintain a sorted list

STL: **vector**

- Constructing/Initializing a **vector**:

```
vector<int> myVector(15);  
vector<double> myReals(20, 137.0); //Initial value 137.0  
vector<string> myStrings(5, "(none)");
```

- Resizing a **vector**:

```
vector<int> myVector; // Defaults to empty vector  
PrintVector(myVector); // Output: [nothing]  
myVector.resize(10); // Grow the vector, setting new elements to 0  
PrintVector(myVector); // Output: 0 0 0 0 0 0 0 0 0 0  
myVector.resize(5); // Shrink the vector  
PrintVector(myVector); // Output: 0 0 0 0 0
```


STL: `vector` (Summary)

API	Description
Constructor: <code>vector<T> ()</code>	<code>vector<int> myVector;</code> Constructs an empty vector.
Constructor: <code>vector<T> (size_type size)</code>	<code>vector<int> myVector(10);</code> Constructs a vector of the specified size where all elements use their default values (for integral types, this is zero).
Constructor: <code>vector<T> (size_type size, const T& default)</code>	<code>vector<string> myVector(5, "blank");</code> Constructs a vector of the specified size where each element is equal to the specified default value.
<code>size_type size() const;</code>	<code>for(int i = 0; i < myVector.size(); ++i) { ... }</code> Returns the number of elements in the vector.
<code>bool empty() const;</code>	<code>while(!myVector.empty()) { ... }</code> Returns whether the vector is empty.
<code>void clear();</code>	<code>myVector.clear();</code> Erases all the elements in the vector and sets the size to zero.

STL : `vector` (Summary)

API	Description
<pre>T& operator [] (size_type position); const T& operator [] (size_type position) const; T& at(size_type position); const T& at(size_type position) const;</pre>	<pre>myVector[0] = 100; int x = myVector[0]; myVector.at(0) = 100; int x = myVector.at(0);</pre> <p>Returns a reference to the element at the specified position. The bracket notation <code>[]</code> does not do any bounds checking and has undefined behavior past the end of the data. The <code>at</code> member function will throw an exception if you try to access data beyond the end.</p>
<pre>void resize(size_type newSize); void resize(size_type newSize, T fill);</pre>	<pre>myVector.resize(10); myVector.resize(10, "default");</pre> <p>Resizes the vector so that it's guaranteed to be the specified size. In the second version, the <code>vector</code> elements are initialized to the value specified by the second parameter. Elements are added to and removed from the end of the vector, so you can't use <code>resize</code> to add elements to or remove elements from the start of the vector.</p>
<pre>void push_back();</pre>	<pre>myVector.push_back(100);</pre> <p>Appends an element to the vector.</p>
<pre>T& back(); const T& back() const;</pre>	<pre>myVector.back() = 5; int lastElem = myVector.back();</pre> <p>Returns a reference to the last element in the vector.</p>

STL : `vector` (Summary)

<pre>T& front(); const T& front() const;</pre>	<pre>myVector.front() = 0; int firstElem = myVector.front();</pre> <p>Returns a reference to the first element in the vector.</p>
<pre>void pop_back();</pre>	<pre>myVector.pop_back();</pre> <p>Removes the last element from the vector.</p>
<pre>iterator begin(); const_iterator begin() const;</pre>	<pre>vector<int>::iterator itr = myVector.begin();</pre> <p>Returns an iterator that points to the first element in the vector.</p>
<pre>iterator end(); const_iterator end() const;</pre>	<pre>while(itr != myVector.end());</pre> <p>Returns an iterator to the element <i>after</i> the last. The iterator returned by end does not point to an element in the vector.</p>
<pre>iterator insert(iterator position, const T& value); void insert(iterator start, size_type numCopies, const T& value);</pre>	<pre>myVector.insert(myVector.begin() + 4, "Hello"); myVector.insert(myVector.begin(), 2, "Yo!");</pre> <p>The first version inserts the specified value into the vector, and the second inserts <code>numCopies</code> copies of the value into the vector. Both calls invalidate all outstanding iterators for the vector.</p>

STL : `vector` (Summary)


```
iterator erase(iterator position);  
iterator erase(iterator start,  
iterator end);
```

```
myVector.erase(myVector.begin());  
myVector.erase(startItr, endItr);
```

The first version erases the element at the position pointed to by position. The second version erases all elements in the range [startItr, endItr). Note that this does **not** erase the element pointed to by endItr. All iterators after the remove point are invalidated

More STL containers

- **deque** : Supports all vector functionality and **push_front()** and **pop_front()**



- **stack** and **queue**

STL: Associative Containers

- **`std::set`**

- An unordered collection of unique elements that does not permit duplicates. Unlike vector and deque, sets can only store objects for which the “<” operator is defined.
- Internally, a set is layered on top of a balanced binary search tree. But since a binary search tree orders elements in a certain order, “<” needs to be defined.

Traversing Containers with Iterators

- What exactly is an iterator? At a high level, it is like a cursor in the text editor.

```
vector<int> myVector = /* ... some initialization ... */  
for (vector<int>::iterator itr = myVector.begin();  
itr != myVector.end(); ++itr)  
cout << *itr << endl;
```

- Every STL container class exports a member function **begin()** which yields an iterator pointing to the first element of that container.
- The strange-looking entity ***itr** is known as an *iterator dereference* and means “the element being iterated over by **itr**.”
- When applied to iterators, the **++** operator means “advance the iterator one step forward.”
- To detect when an iterator has visited all of the elements, we loop on the condition that

```
itr != myVector.end();
```

Traversing Containers with Iterators

- Notice that the **begin()** iterator points to the first element of the vector, while the **end()** iterator points to the slot one position past the end of the container.

`std::pair<type1, type2>`

- Make a **pair**:

```
pair<int, string> myPair;  
myPair.first = 137;  
myPair.second = "C++ is awesome!";  
  
pair<int, string> myPair = make_pair(137, "string!");
```

`std::map< KeyType, ValueType >`

- **`std::map`**

- Maps are associative containers that allow you to store data indexed by keys.
- Maps can be accessed using iterators that are essentially pointers to templated objects of base type `pair`, which has two members, first and second. First corresponds to the key, second to the value associated with the key.
- Maps are fast, guaranteeing $O(\log(n))$ insertion and lookup time. (Think BST)
- To check if something is not in the **`map`**:

```
std::map <string, char> grade_list;
grade_list["John"] = 'A';
if(grade_list.find("Tim") == grade_list.end())
{
    std::cout<<"Tim is not in the map!"<<endl;
}
```

`std::map< KeyType, ValueType >`

- Because the square brackets both query and create key/value pairs, you should use care when looking values up with square brackets.
- If you want to look up a key/value pair without accidentally adding a new key/value pair to the **map**, you can use the **map**'s **find** member function. **find** takes in a key, then returns an iterator that points to the key/value pair that has the specified key. If the key does not exist, **find** returns the **map**'s **end()** iterator.
- **map** iterators are slightly more complicated because they dereference to a key/value pair. In particular, if you have a **map<KeyType, ValueType>**, then the iterator will dereference to a value of type **pair<const KeyType, ValueType>**

`std::unordered_map<keyType, valueType>`

- Unordered maps are associative containers that store elements formed by the combination of a *key value* and a *mapped value*, and which allows for fast retrieval of individual elements based on their keys. Hash Maps!!!

STL Algorithms

- Let's say you want to find the average of the values in a **vector**, you can loop over, perform the summation and divide by number of elements or simply do:

```
accumulate(values.begin(), values.end(), 0.0) / values.size()
```

- The power comes in its ability to handle a wide range of STL containers like, **vector**, **deque**, **set** etc.
- Reasons to do this: simplicity, clarity and correctness

STL Algorithms

- The suffix `_if` on an algorithm (**replace_if**, **count_if**, etc.) means the algorithm will perform a task on elements only if they meet a certain criterion.

```
count(myVec.begin(), myVec.end(), 137)
```

```
bool IsEven(int value)
{
    return value % 2 == 0;
}

cout << count_if(myVec.begin(), myVec.end(), IsEven) << endl;
```

- Copy and fill algorithms can be use to copy a range/initialize (clear) a range.

STL Algorithms

- Reordering Algorithms:

```
bool CompareStringLength(string one, string two)
{
    return one.length() < two.length();
}
sort(myVector.begin(), myVector.end(), CompareStringLength);
```

- See <https://en.cppreference.com/w/cpp/algorithm> for the full list

References

- The C++ Programming Language, 4th Edition, Bjarne Stroustrup
- Thinking in C++, Bruce Eckel
- Effective Modern C++ , Scott Meyers
- Quick Guide to Modern C++: A Tour of C++, 2nd edition

Course Survey

- When you have 2 minutes, please fill out:
- <https://www.surveymonkey.com/r/XKSGTM3>